

Building a Serverless Web Application on AWS Services with Python: A Comprehensive Guide for Student Data Management

This document provides a detailed, step-by-step guide for constructing a modern serverless web application on Amazon Web Services (AWS) using Python for the backend logic. This architecture emphasizes scalability, cost-effectiveness, and operational efficiency, allowing developers to concentrate on application features rather than infrastructure management.

Project Overview

This project outlines the creation of a serverless web application designed to manage student data. By leveraging AWS's fully managed serverless offerings, we aim to demonstrate a robust and scalable solution for data storage, retrieval, and presentation without the overhead of traditional server provisioning. The choice of Python for AWS Lambda functions ensures a widely adopted and versatile language for backend processing.

Step-by-Step Implementation Guide

Step 1: Backend Setup (Amazon DynamoDB and AWS Lambda with Python)

This section details the configuration of the core backend components responsible for data storage and business logic.

1.1. Create an Amazon DynamoDB Table

The DynamoDB table will serve as the persistent storage for student records.

1. **Navigate to DynamoDB:** Access the AWS Management Console and open the DynamoDB service.
2. **Initiate Table Creation:** Click "Create table".---(Figure 2)

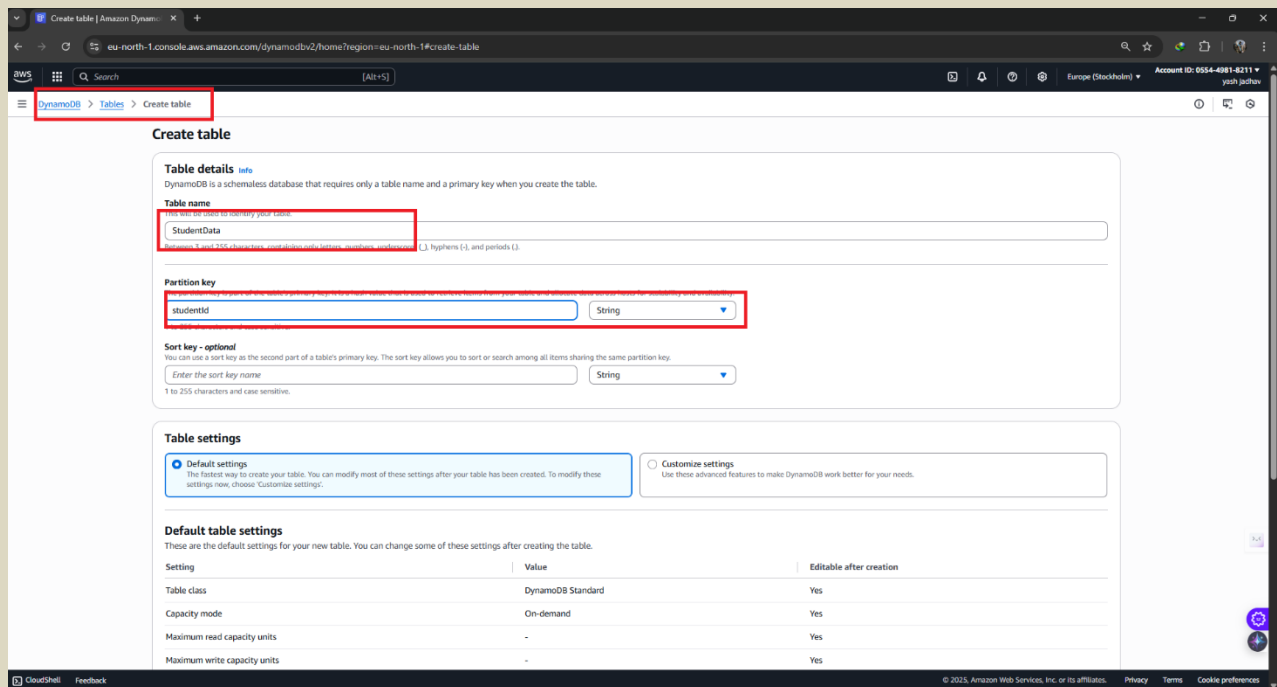


Figure 2 Create DynamoDB

3. Define Table Properties:

- **Table name:** Enter a descriptive name, such as StudentData.
- **Partition key:** Specify studentId (String type) as the primary key. This uniquely identifies each student record. --- (Figure 3)
- **Sort key (Optional):** For more complex queries, a Sort Key can be added (e.g., courseId if students can enroll in multiple courses, and you want to retrieve courses for a specific student efficiently). For this guide, it's optional.
- **Default settings:** Retain the default settings for initial setup.

4. Finalize Table Creation: Click "Create table".

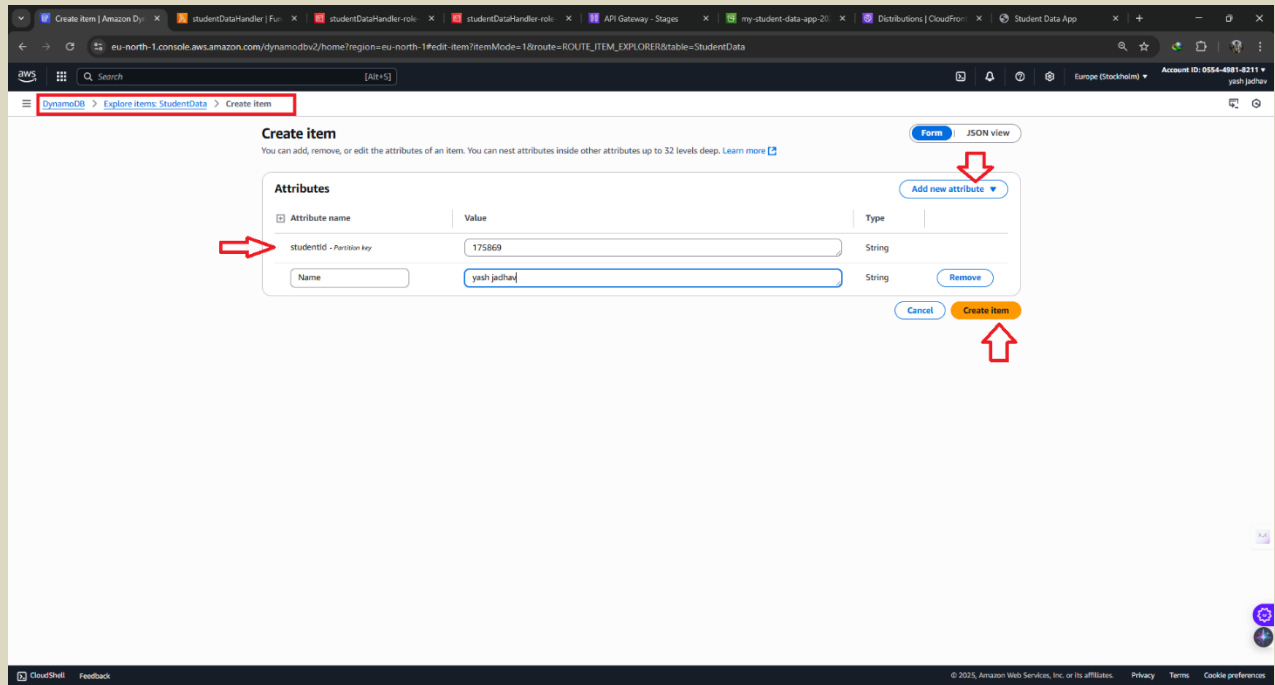


Figure 3 Add table and Create item in DynamoDB

1.2. Create an AWS Lambda Function

The Lambda function will execute the Python code for interacting with DynamoDB.

1. **Navigate to Lambda:** Access the AWS Management Console and open the Lambda service.
2. **Initiate Function Creation:** Click "Create function".
3. **Configure Basic Information:**
 - **Authoring option:** Select "Author from scratch".
 - **Function name:** Provide a clear name, e.g., studentDataHandler.
 - **Runtime:** Choose "Python 3.9" (or a newer supported version) to execute your Python code.
 - **Architecture:** Select x86_64 for general compatibility.
 - **Execution role:** Under "Change default execution role", you can create a new role or select an existing one. A new role typically gets Basic Lambda permissions which includes CloudWatch Logs permissions.
4. **Finalize Function Creation:** Click "Create function". ---(Figure 4)

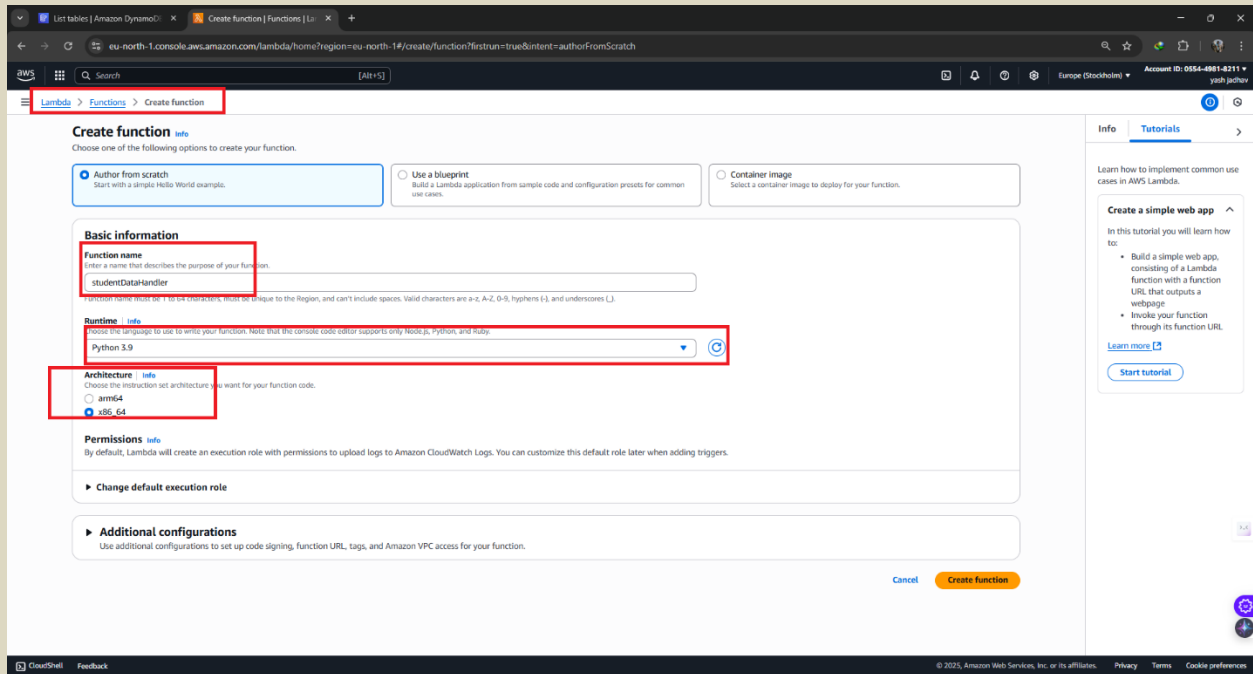
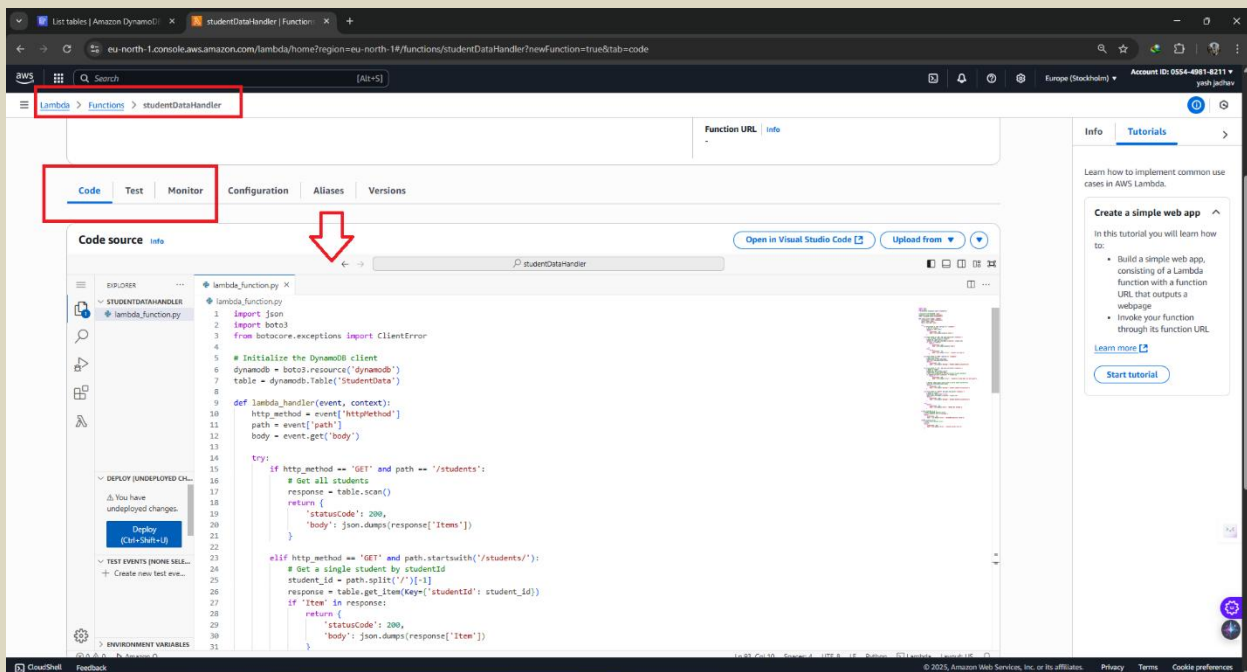


Figure 4 Create a Lambda Function

1.3. Develop Python Lambda Code for CRUD Operations

Replace the default lambda handler code in the Lambda console's code editor with the following Python code. This function processes HTTP requests from API Gateway to perform CRUD operations on the StudentData DynamoDB table. Copy the code in source code of lambda function ---(Figure below)



```
import json
import boto3
from botocore.exceptions import ClientError

# Initialize the DynamoDB client
dynamodb = boto3.resource('dynamodb')
table = dynamodb.Table('StudentData')

def lambda_handler(event, context):
    http_method = event['httpMethod']
    path = event['path']
    body = event.get('body')

    # Define a standard response dictionary with headers
    response_headers = {
        'Content-Type': 'application/json',
        'Access-Control-Allow-Origin': '*'
    }

    try:
        if http_method == 'GET' and path == '/students':
            # Get all students
            response = table.scan()
            return {
                'statusCode': 200,
                'headers': response_headers,
                'body': json.dumps(response['Items'])
            }

        elif http_method == 'GET' and path.startswith('/students/'):
            # Get a single student by studentId
            student_id = path.split('/')[-1]
            response = table.get_item(Key={'studentId': student_id})
            if 'Item' in response:
                return {
                    'statusCode': 200,
                    'headers': response_headers,
                    'body': json.dumps(response['Item'])
                }
            else:
                return {
                    'statusCode': 404,
                    'headers': response_headers,
                    'body': json.dumps({'error': 'Student not found'})
                }

        elif http_method == 'POST' and path == '/students':
            # Add a new student
```

```

student_data = json.loads(body)
table.put_item(Item=student_data)
return {
    'statusCode': 201,
    'headers': response_headers,
    'body': json.dumps({'message': 'Student added successfully'})
}

elif http_method == 'PUT' and path.startswith('/students/'):
    # Update a student
    student_id = path.split('/')[-1]
    student_data = json.loads(body)
    if student_data.get('studentId') != student_id:
        return {
            'statusCode': 400,
            'headers': response_headers,
            'body': json.dumps({'error': 'studentId in body does not match path'})
        }

    table.put_item(Item=student_data)
    return {
        'statusCode': 200,
        'headers': response_headers,
        'body': json.dumps({'message': 'Student updated successfully'})
    }

elif http_method == 'DELETE' and path.startswith('/students/'):
    # Delete a student
    student_id = path.split('/')[-1]
    table.delete_item(Key={'studentId': student_id})
    return {
        'statusCode': 200,
        'headers': response_headers,
        'body': json.dumps({'message': 'Student deleted successfully'})
    }

else:
    return {
        'statusCode': 405,
        'headers': response_headers,
        'body': json.dumps({'error': 'Method Not Allowed'})
    }

except ClientError as e:
    print(e.response['Error']['Message'])
    return {
        'statusCode': 500,
        'headers': response_headers,

```

```

        'body': json.dumps({'error': 'DynamoDB operation failed'})
    }
except Exception as e:
    print(e)
    return {
        'statusCode': 500,
        'headers': response_headers,
        'body': json.dumps({'error': 'Internal server error'})
    }

```

1.4. Configure IAM Permissions for AWS Lambda

The Lambda function requires appropriate permissions to interact with DynamoDB and log to CloudWatch.

1. **Access Lambda Function Permissions:** Navigate to the studentDataHandler function in the Lambda console, then click the **"Configuration"** tab and select **"Permissions"**.— (Figure 5)
2. **Edit Execution Role:** Click on the **Role name** (e.g., studentDataHandler-role-xxxxxx) to open the IAM console for that role.

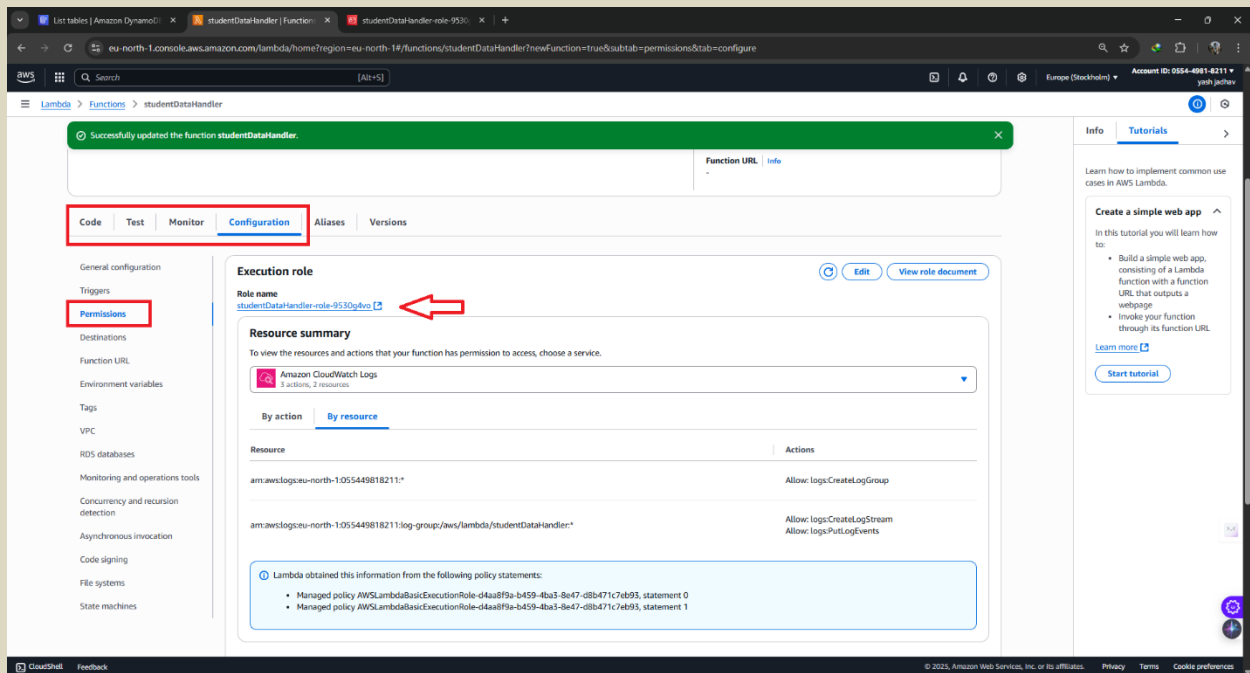


Figure 5 Lambda Function Permissions

3. **Add Inline Policy:**---(Figure 6)
 - o In the IAM role, click **"Add permissions"** -> **"Create inline policy"**.
 - o **Service:** Choose DynamoDB.
 - o **Actions:** Under "Access level", expand Write, Read, and Permissions management.

Select the following actions:

- GetItem
- PutItem
- DeleteItem
- Scan
- UpdateItem (although put_item is used in the example for updates, UpdateItem is more granular for partial updates and good to include for future expansion).

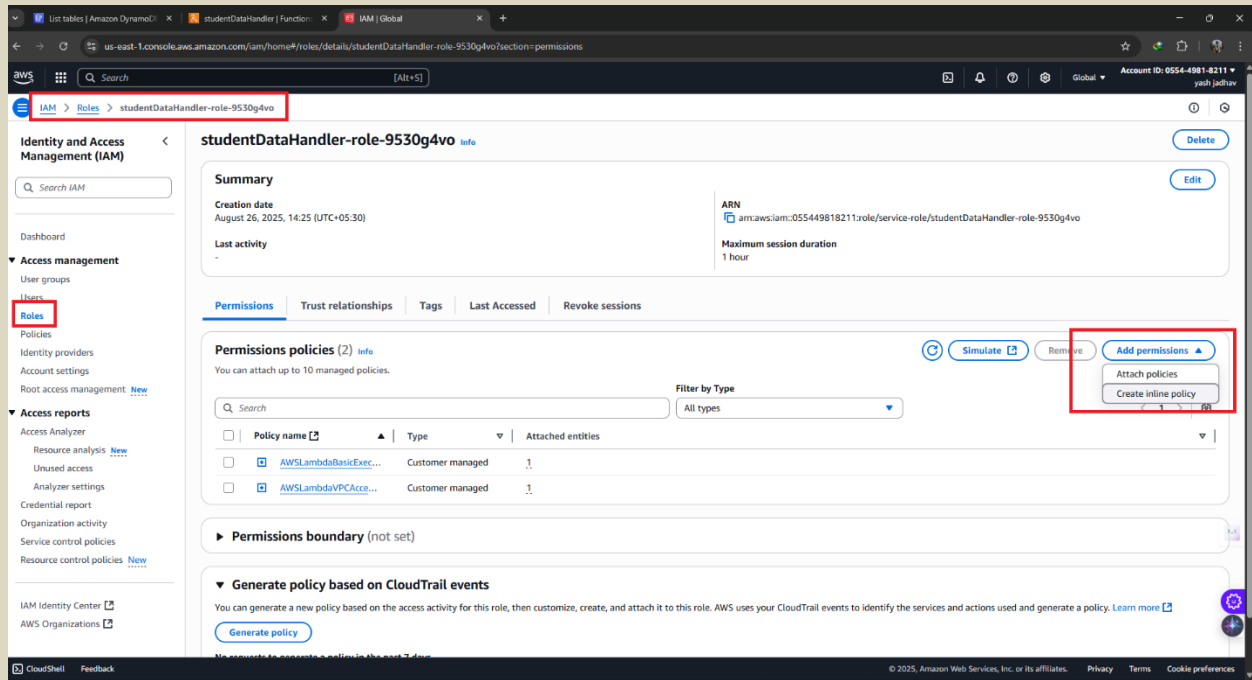


Figure 7 Create inline policy

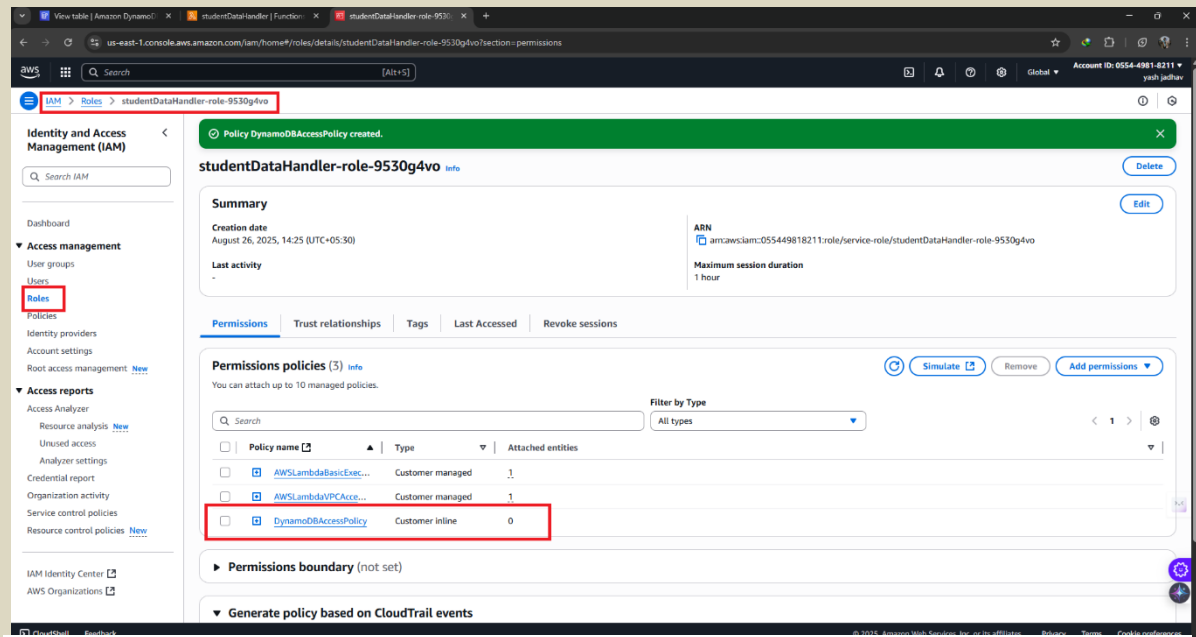
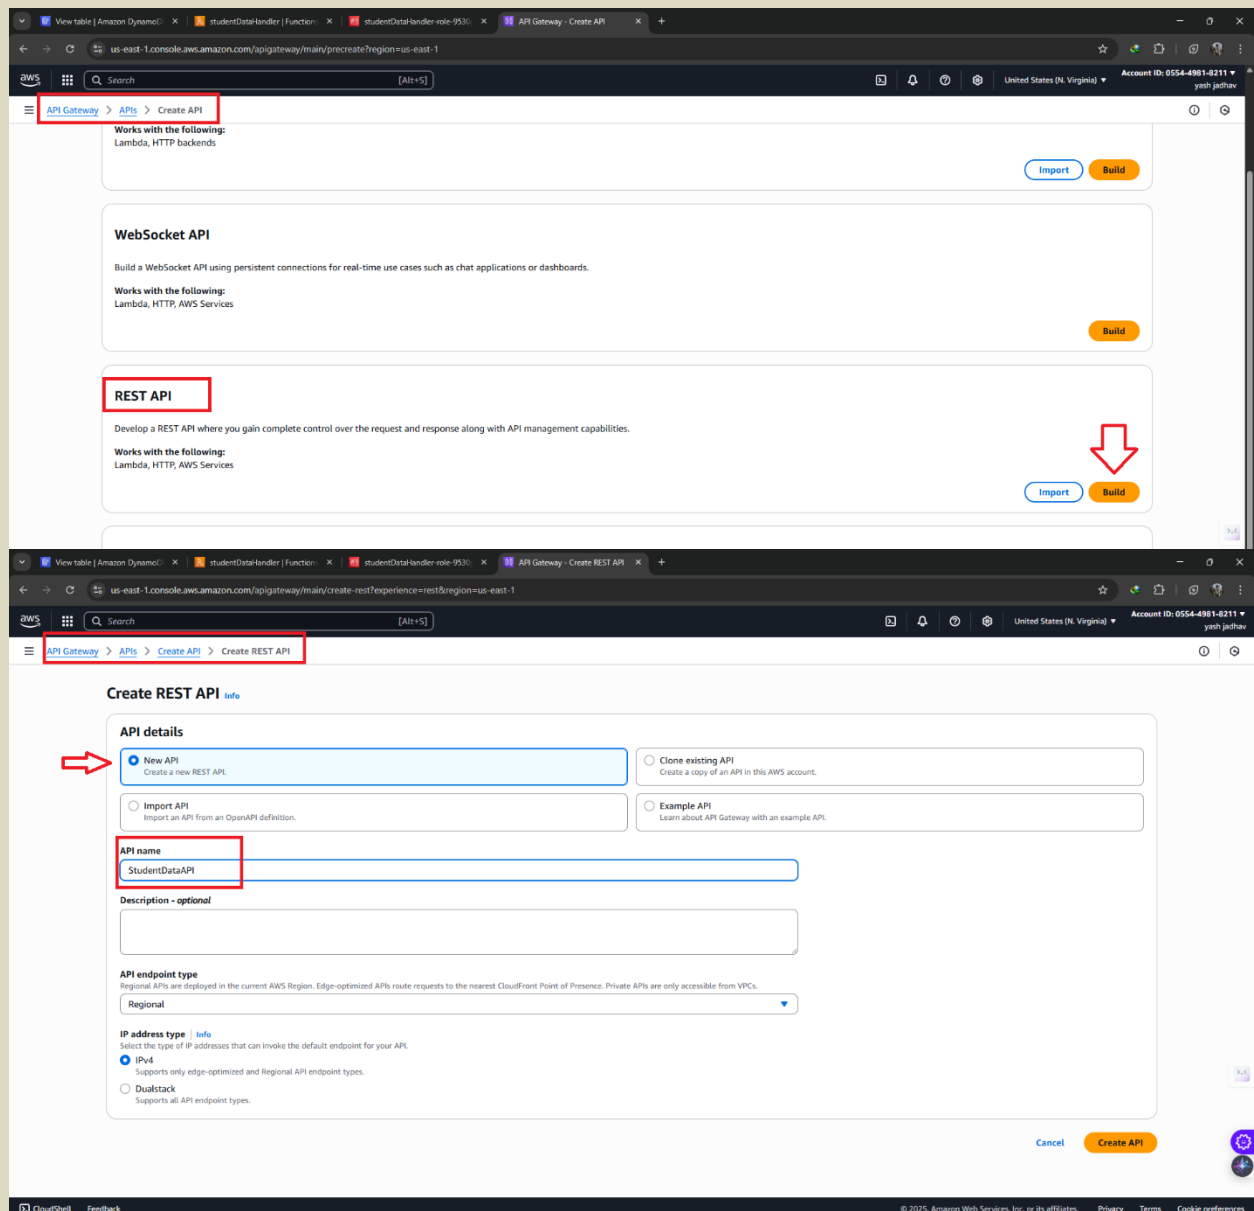


Figure 6 created a policy

Step 2: Set up Amazon API Gateway

API Gateway provides the public HTTP endpoint for clients to interact with the Lambda backend.

1. **Navigate to API Gateway:** Access the AWS Management Console and open the API Gateway service.
2. **Create a New REST API:**
 - Click "Build" under "REST API".
 - Choose "New API".---(Figure 8)
 - **API name:** Enter a name, e.g., StudentDataAPI.
 - **Endpoint Type:** Select Regional.
 - Click "Create API".



2.1. Create Resources

Resources define the paths for your API endpoints.

1. Create /students Resource:

- Select the root (/) resource.
- Click "Actions" -> "Create Resource".
- **Resource Name:** Type students.
- **Resource Path:** It will auto-populate as /students.
- Click "Create Resource".

2. Create /students/{studentId} Resource: This path will be used for operations on individual student records.

- Select the /students resource you just created.
- Click "Actions" -> "Create Resource".
- **Resource Name:** Type studentId.
- **Resource Path:** It will auto-populate as /{studentId}.
- Check "Use Lambda Proxy integration".
- Click "Create Resource".

2.2. Create Methods for Resources

Methods (GET, POST, PUT, DELETE) define the HTTP operations for your resources. We will use Lambda Proxy Integration for simplicity and flexibility.

1. For /students Resource:

- Select the /students resource.
- Click "Actions" -> "Create Method".
- Choose GET.
- **Integration type:** Select Lambda Function.
- Check "Use Lambda Proxy integration".
- **Lambda Region:** Select the region where your Lambda function resides.
- **Lambda Function:** Enter studentDataHandler and click the checkmark.
- A pop-up will appear asking to add permission to Lambda function. Click "OK".
- Repeat the process for POST method on /students resource, integrating it with studentDataHandler.

2. For /students/{studentId} Resource:

- Select the /students/{studentId} resource.
- Click "Actions" -> "Create Method".
- Choose GET. Configure with studentDataHandler using Lambda Proxy Integration.
- Repeat for PUT and DELETE methods, also integrating with studentDataHandler using Lambda Proxy Integration.---(Figure 8)

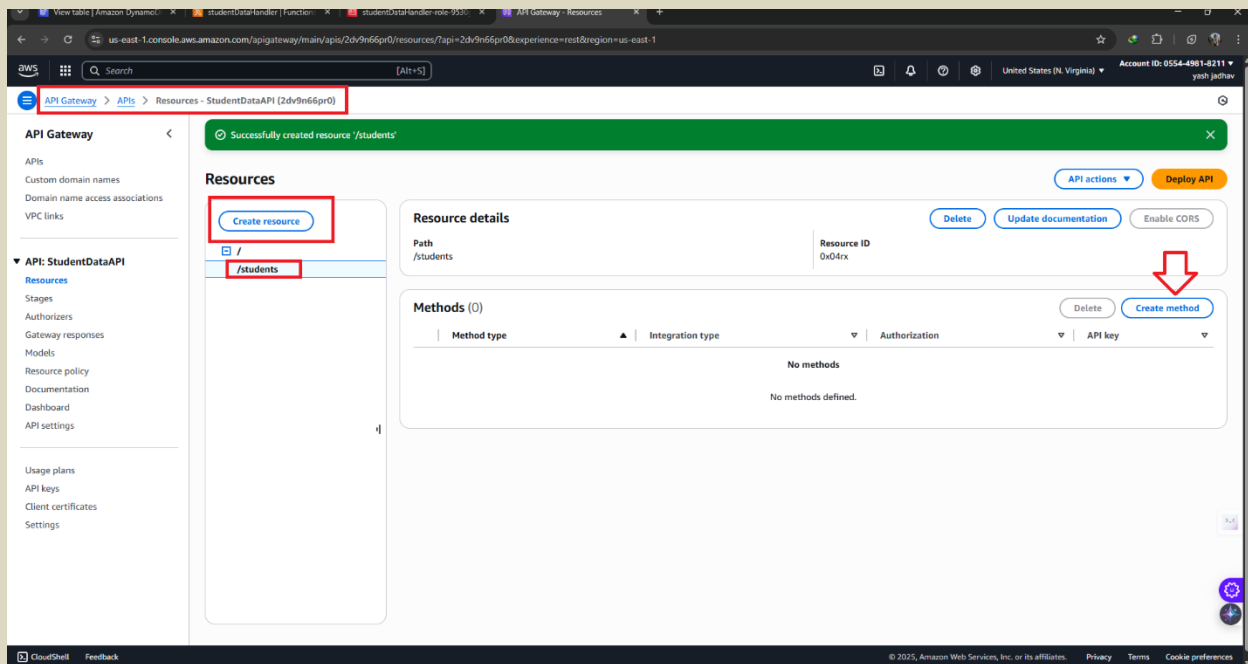


Figure 8 API RESOURCE

2.3. Enable CORS for the API

For a web application hosted on a different domain (S3/CloudFront), Cross-Origin Resource Sharing (CORS) must be enabled.

1. Select the /students resource (and if needed, the root / resource and /students/{studentId} resource).
2. Click "Actions" -> "Enable CORS".
3. For a simple setup, you can keep the default settings (allowing GET,HEAD,POST,PUT,DELETE,OPTIONS from * origin). In a production environment, restrict the Access-Control-Allow-Origin to your CloudFront domain.
4. Click "Enable CORS and replace existing CORS headers".

2.4. Deploy the API

Deployment makes your API accessible via a public URL.

1. Click "Actions" -> "Deploy API".
2. **Deployment stage:** Select "[New Stage]".
3. **Stage name:** Give it a name, e.g., prod.
4. Click "Deploy".
5. After deployment, note the **Invoke URL**. This URL, combined with your resource paths (e.g., https://xxxxxx.execute-api.region.amazonaws.com/prod/students), will be used by your frontend.---(Figure 9)

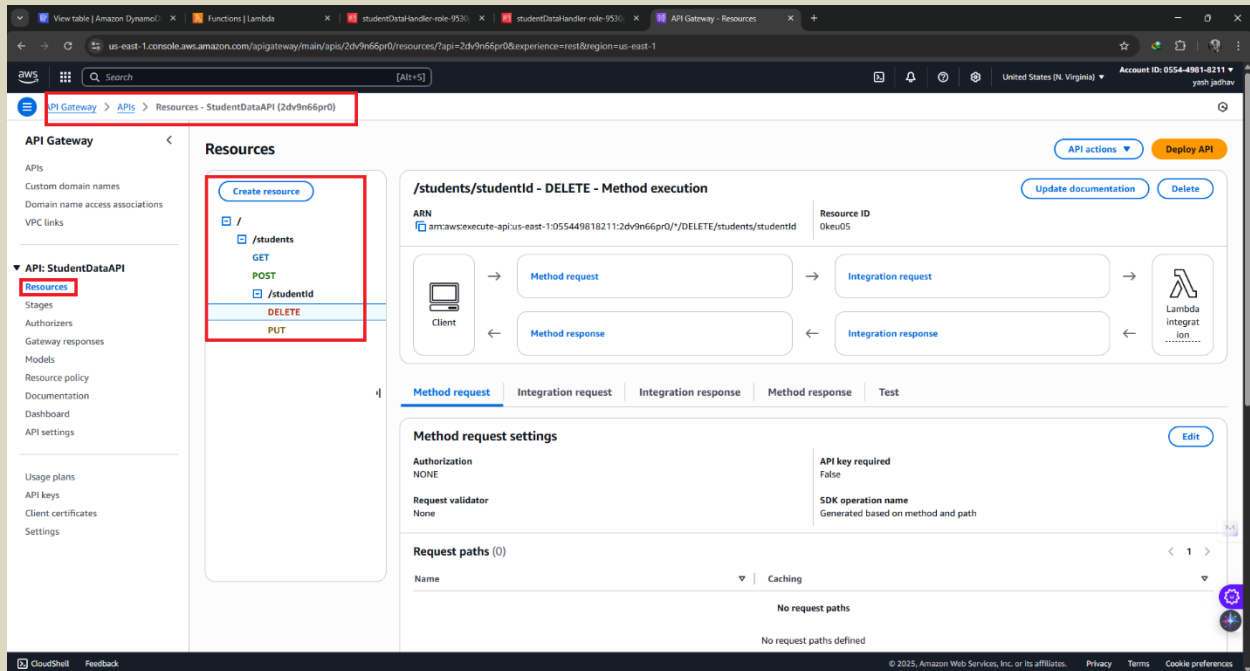


Figure 9 Methods and Deploy

Step 3: Frontend Setup (Amazon S3 and Amazon CloudFront)

This section covers hosting the static frontend assets and distributing them globally.

3.1. Create an Amazon S3 Bucket for Static Website Hosting

The S3 bucket will store your HTML, CSS, and JavaScript files.

1. **Navigate to S3:** Access the AWS Management Console and open the S3 service.
2. **Initiate Bucket Creation:** Click "Create bucket".
3. **Configure Bucket Properties:**
 - **Bucket name:** Choose a globally unique name, e.g., my-student-data-app-2025-frontend.
 - **AWS Region:** Select a region geographically close to your users or primary backend services.
 - **Object Ownership:** Keep ACLs disabled (recommended).
 - **Block Public Access settings for this bucket:** Uncheck "Block all public access" and acknowledge the warning. This is crucial for enabling public access required by static website hosting. *Note: For a more secure approach, especially when using CloudFront, you would typically keep S3 private and use an Origin Access Control (OAC) with CloudFront to restrict direct S3 access.*
4. **Finalize Bucket Creation:** Click "Create bucket".—(Figure 10)

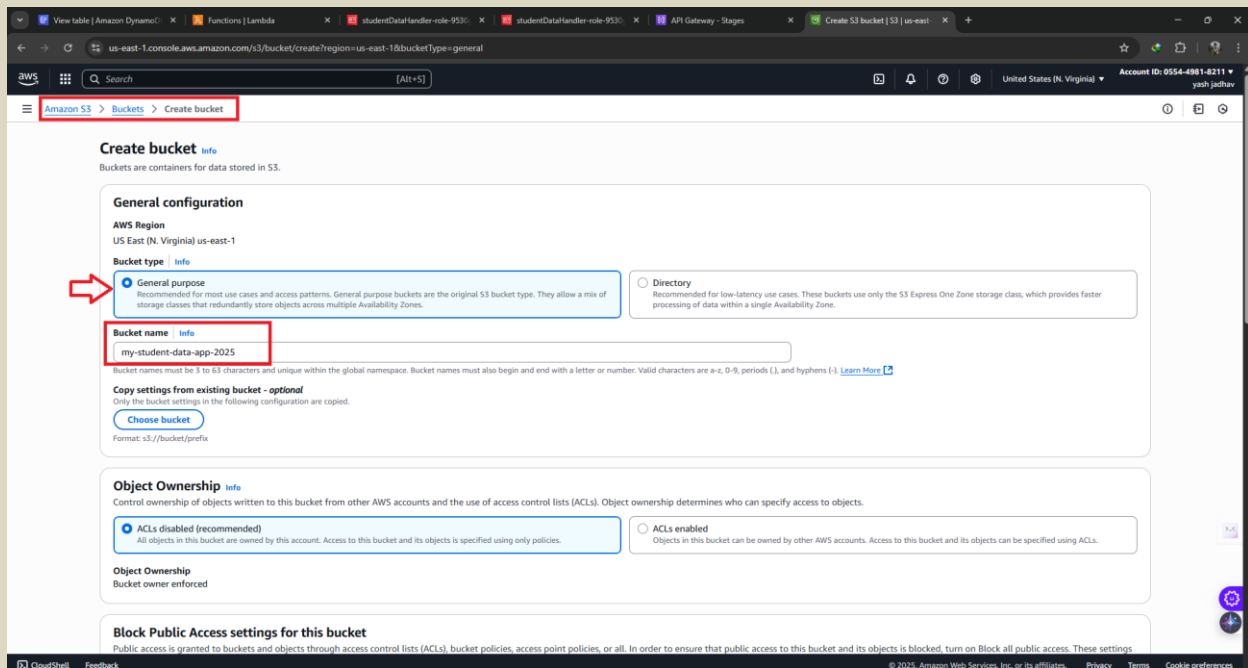


Figure 10 S3 bucket

3.2. Enable Static Website Hosting on the S3 Bucket

1. **Access Bucket Properties:** Select your newly created S3 bucket and go to the "**Properties**" tab.
2. **Configure Static Website Hosting:**
 - Scroll down to "**Static website hosting**" and click "**Edit**".
 - Select "**Enable**".
 - **Index document:** Enter index.html.—(Figure 11)
 - **Error document (Optional):** Enter error.html for custom error pages.
3. **Save Changes:** Click "**Save changes**". Note the **Bucket website endpoint** from this section (e.g., <http://my-student-data-app-2025-frontend.s3-website.region.amazonaws.com>).

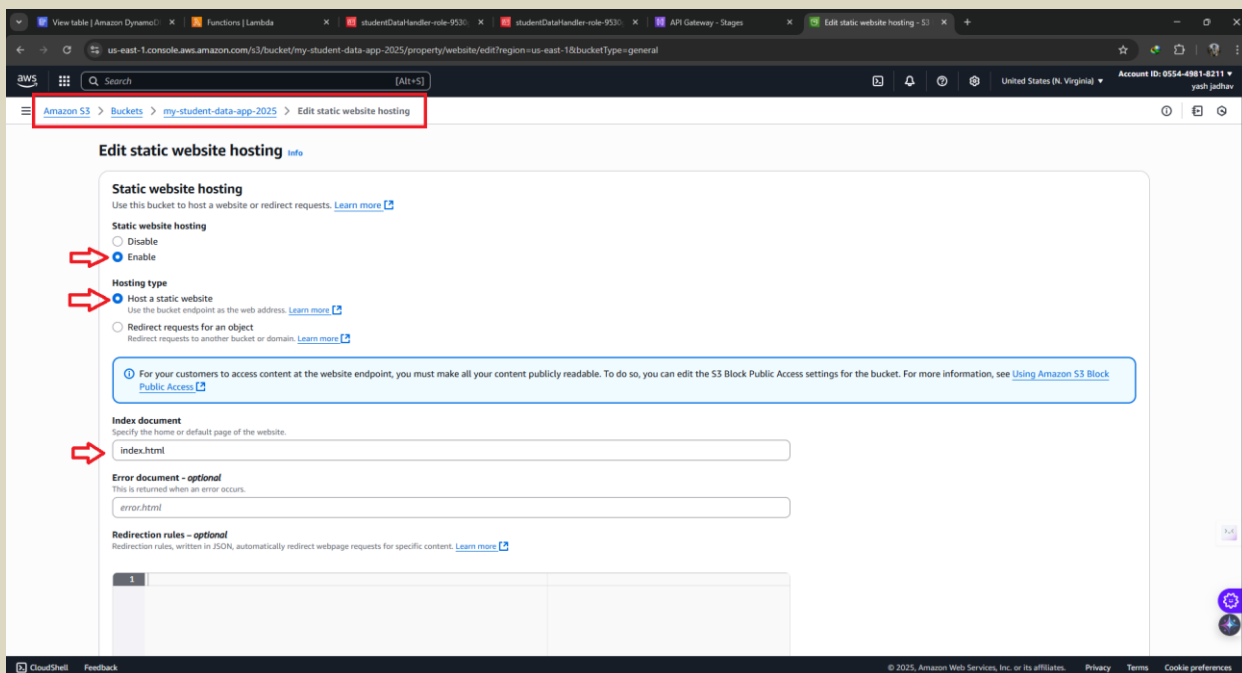


Figure 11 Enable Static Website Hosting on the S3 Bucket

3.3. Create and Upload Frontend Assets

Create your static web files (index.html, styles.css, script.js if separate) and upload them to the S3 bucket.

***index.html (Example):**

```
<!DOCTYPE html>
<link rel="stylesheet" href="style.css">
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Student Data App</title>
  <link rel="stylesheet" href="style.css">
</head>
<body>
  <h1>Student Data Management</h1>
  <div id="student-list"></div>
  <script>
    const API_URL = 'https://bsf6c7ib6g.execute-api.eu-north-
1.amazonaws.com/prod/students';{YOUR API}

    async function fetchStudents() {
      const response = await fetch(API_URL);
      const students = await response.json();
      const listElement = document.getElementById('student-list');
      listElement.innerHTML = ""; // Clear previous data
      students.forEach(student => {
        const studentDiv = document.createElement('div');
        studentDiv.textContent = `ID: ${student.studentId}, Name: ${student.name}, Percentage:
${student.percentage}`;
        listElement.appendChild(studentDiv);
      });
    }

    fetchStudents();
  </script>
</body>
</html>
```


***styles.css (Example):**

```
/* General body and typography */
body {
  font-family: -apple-system, BlinkMacSystemFont, 'Segoe UI', Roboto, Helvetica, Arial, sans-serif;
  background-color: #f4f7f6;
  color: #333;
  line-height: 1.6;
  margin: 0;
  padding: 20px;
  display: flex;
  flex-direction: column;
  align-items: center;
}

/* Header styling */
h1 {
  color: #2c3e50;
  text-align: center;
  margin-bottom: 30px;
  font-weight: 600;
  letter-spacing: 1px;
  border-bottom: 2px solid #3498db;
  padding-bottom: 10px;
}

/* Student list container */
#student-list {
  width: 100%;
  max-width: 600px;
  background-color: #ffffff;
  border-radius: 8px;
  box-shadow: 0 4px 12px rgba(0, 0, 0, 0.05);
  padding: 20px;
}

/* Individual student item styling */
#student-list div {
  padding: 15px;
  border-bottom: 1px solid #e9ecef;
  transition: background-color 0.3s ease;
  font-size: 1rem;
  color: #555;
}

#student-list div:last-child {
  border-bottom: none;
}

#student-list div:hover {
```

```

    background-color: #e8f4f8;
}

/* Responsive design */
@media (max-width: 768px) {
    body {
        padding: 10px;
    }

    h1 {
        font-size: 1.5rem;
    }

    #student-list {
        padding: 10px;
    }
}

```

1. **Upload Assets:** Upload index.html and styles.css to your S3 bucket. Ensure the API_BASE_URL in index.html is updated with your API Gateway Invoke URL.

3.4. Create an Amazon CloudFront Distribution

CloudFront will serve your static content securely and efficiently.

1. **Navigate to CloudFront:** Access the AWS Management Console and open the CloudFront service.
2. **Initiate Distribution Creation:** Click "Create Distribution".
3. **Configure Origin Settings:**
 - o **Origin domain:** Select your S3 bucket's **Bucket website endpoint** from the dropdown list (e.g., http://my-student-data-app-2025-frontend.s3-website.region.amazonaws.com). *Do NOT select the S3 REST API endpoint that CloudFront automatically suggests if you want S3 static website hosting features like index documents.*
 - o CloudFront will automatically populate the **Origin ID**.—(Figure 13)
4. **Default Cache Behavior Settings:**
 - o **Viewer protocol policy:** Select Redirect HTTP to HTTPS for secure access.
 - o **Allowed HTTP methods:** Ensure GET, HEAD, OPTIONS are allowed (or more if your frontend uses other methods for static assets).
 - o **Cache policy:** Use CachingOptimized or create a custom policy.
5. **Distribution Settings:**
 - o **Default root object:** Enter index.html. This ensures that when users access your domain, index.html is served by default.—(Figure 14)
 - o **Price class:** Choose based on your geographic reach and cost preferences.
 - o **Web Application Firewall (WAF) (Optional):** Consider enabling AWS WAF for additional protection against common web exploits.

6. **Finalize Distribution Creation:** Click "Create Distribution". This process can take several minutes to complete.
7. **Distribution Domain Name:** Once deployed, note the **Distribution domain name** (e.g., d1234abcd.cloudfront.net). This is the secure URL for your web application.—(Figure 14)

The screenshot shows the AWS CloudFront 'Create distribution' console page. The 'Distribution name' field is highlighted with a red box and contains the text 'Cloud-front-for-s3'. The 'Distribution type' section shows 'Single website or app' selected. The 'Custom domain' section is empty. The 'Tags' section is also empty. A 'Next' button is visible at the bottom right.

Figure 14 Add s3 and root object as index.html

The screenshot shows the AWS CloudFront console details for a newly created distribution named 'Cloud-front-for-s3'. The 'Distribution domain name' is highlighted with a red box and shows 'dfeq3c78wqy.cloudfront.net'. The 'Settings' section shows 'Standard logging' and 'Default root object' as '-'. The 'Continuous deployment' section has a 'Create staging distribution' button.

Figure 13 Add your S3 ,API in cloud front and change root object too index.html

3.5. Add Data to your Dynamodb and Paste your ARN in Browsers

Create item

You can add, remove, or edit the attributes of an item. You can nest attributes inside other attributes up to 32 levels deep. [Learn more](#)

Attributes

Attribute name	Value	Type
studentId - Partition key	175969	String
Name	yash jadhav	String

[Add new attribute](#) [Remove](#) [Cancel](#) [Create item](#)

Student Data Management

- ID: 142596, Name: yash jadhav
- ID: 1456789, Name: soham
- ID: yash jadhav, Name: undefined

Enhanced Security and Best Practices

Monitoring and Logging

- **AWS CloudWatch:** All AWS services used (Lambda, API Gateway, DynamoDB) integrate with CloudWatch for logging and metrics.
 - **Lambda Logs:** Lambda automatically sends function logs to CloudWatch Logs. Configure custom log groups and metrics for advanced analysis.
 - **API Gateway Logs:** Enable CloudWatch Logs for API Gateway to capture detailed request and response information, including errors and access logs.
 - **CloudWatch Alarms:** Set up alarms on critical metrics (e.g., Lambda errors, API Gateway 5xx errors, DynamoDB throttled events) to notify administrators of operational issues.
- **AWS X-Ray (Distributed Tracing):** Integrate AWS X-Ray with Lambda and API Gateway to trace requests end-to-end. This helps visualize application components, identify performance bottlenecks, and debug issues across services.

Cost Optimization

- **Serverless Pricing Model:** Pay only for the compute time and resources consumed. Lambda charges per invocation and duration, DynamoDB charges for read/write capacity units (RCUs/WCUs) or on-demand, and API Gateway charges per million API calls.
- **DynamoDB On-Demand vs. Provisioned:** For unpredictable workloads, DynamoDB On-Demand capacity can be more cost-effective. For stable or predictable workloads, Provisioned capacity with Auto Scaling can optimize costs.
- **Lambda Memory Allocation:** Optimize Lambda function memory. Higher memory often correlates with more CPU, so finding the right balance is crucial for performance and cost.
- **CloudFront Cache Policy:** Optimize caching policies to minimize requests to your S3 origin, reducing data transfer costs.

Deployment and CI/CD

- **Infrastructure as Code (IaC):** Utilize IaC tools like AWS Serverless Application Model (SAM) or AWS CloudFormation (or Terraform) to define and manage your entire serverless application stack. This enables version control, automated deployments, and consistent environments.
- **CI/CD Pipeline:** Implement a Continuous Integration/Continuous Deployment (CI/CD) pipeline using services like AWS CodePipeline, CodeBuild, and CodeDeploy. This automates testing, building, and deploying your application changes, enabling faster and more reliable releases.

Frontend Enhancements

The provided index.html is a basic starting point. For a production-ready application, consider:

- **Form for Adding/Updating:** Implement a dedicated form for submitting student data,

including validation.

- **Dynamic UI Updates:** Use JavaScript to dynamically update the student list without full page reloads.
- **Error Handling in UI:** Display user-friendly error messages for API failures.
- **Loading States:** Show loading indicators during API calls.
- **More Student Attributes:** Expand student data (e.g., age, grades, contact information) in both frontend and backend.
- **Styling Libraries:** Integrate modern CSS frameworks (e.g., Tailwind CSS, Bootstrap) for a polished look and responsive design.

This comprehensive guide provides a foundation for building a robust and scalable serverless web application on AWS using Python. By following these steps and incorporating best practices, you can develop efficient and secure student data management solutions.

This video provides an overview of how to securely host a static website with S3 and CloudFront. [Host a Static Website Securely with S3 & CloudFront Step-by-step](#)